

Processes

- **Definition: A *process* is an instance of a running program.**
 - One of the most profound ideas in computer science
 - Not the same as “program” or “processor”
- **Process provides each program with two key abstractions:**
 - Logical control flow
 - Each program seems to have exclusive use of the CPU
 - Private virtual address space
 - Each program seems to have exclusive use of main memory
- **How are these Illusions maintained?**
 - Process executions interleaved (multitasking) or run on separate cores
 - Address spaces managed by virtual memory system

INT 21h – The Gateway to MS-DOS

- In the 16-bit era (MS-DOS), INT 21h was the **primary Application Programming Interface (API)**. It provided all "DOS Services," including file management, console I/O, and program execution control. It acted as a "software gate." Instead of calling a memory address (which might change), the programmer just called INT 21h, and the OS handled the rest. While modern 64-bit **Windows** no longer uses INT 21h, it remains the architectural "ancestor" of how software communicates with a kernel.
- **Parameters for INT 21h** were not passed on a stack (like in C) but directly in CPU Registers. The value in the AH register determined which service was being requested.


```
mov ah, 09h      ; Select "Print String" service
mov dx, offset msg ; Provide the address of the data
int 21h          ; Trigger the software interrupt
```

Modern "Interrupts" – SYSCALL and SYSENTER

- In modern **Windows**, the transition from User Mode to Kernel Mode is handled by specialized hardware instructions: **SYSCALL** (AMD/Intel 64-bit) or **SYSENTER** (Intel 32-bit).
- **The Process:**
 1. Application calls a Win32 function.
 2. The function calls a "Native API" in ntdll.dll.
 3. ntdll.dll puts the system call number in the RAX register and executes SYSCALL.
 4. The CPU immediately jumps to a fixed location in the Kernel.


Concurrent Processes

- Two processes *run concurrently* (are concurrent) if their flows overlap in time
- Otherwise, they are *sequential*
- Examples (running on single core):
 - Concurrent: A & B, A & C
 - Sequential: B & C




User View of Concurrent Processes

- Control flows for concurrent processes are physically disjoint in time
- However, we can think of concurrent processes as running in parallel with each other



Context Switching

- Processes are managed by a shared chunk of OS code called the *kernel*
 - Important: the kernel is not a separate process, but rather runs as part of some user process
- Control flow passes from one process to another via a *context switch*



fork: Creating New Processes

CreateProcess: Windows

■ `int fork(void)`

- creates a new process (child process) that is identical to the calling process (parent process)
- returns 0 to the child process
- returns child's **pid** (process id) to the parent process

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

- Fork is interesting (and often confusing) because it is called *once* but returns *twice*
- *CreateProcess* creates a new, empty address space and loads the specified program into it from disk.

Understanding fork

Process n




```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```


Child Process m



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

hello from parent

Which one is first?

hello from child

The World of Multitasking

- **System runs many processes concurrently**
- **Process: executing program**
 - State includes memory image + register values + program counter
- **Regularly switches from one process to another**
 - Suspend process when it needs I/O resource or timer event occurs
 - Resume process when I/O available or given scheduling priority
- **Appears to user(s) as if all processes executing simultaneously**
 - Even though most systems can only execute one process at a time
 - Except possibly with lower performance than if running alone

Programmer's Model of Multitasking (Linux)


■ Basic functions

- **fork** spawns new process
 - Called once, returns twice
- **exit** terminates own process
 - Called once, never returns
 - Puts it into “zombie” status
- **wait** and **waitpid** wait for and reap terminated children
- **execve** runs new program in existing process
 - Called once, (normally) never returns


■ Programming challenge

- Understanding the nonstandard semantics of the functions
- Avoiding improper use of system resources
 - E.g. “Fork bombs” can disable a system

Unix Process Hierarchy



Windows



Multitasking (Windows)

- **A thread is the smallest unit of execution that can be scheduled by the operating system.**
 - **Process:** A container for resources (address space, file handles, security tokens). A process "owns" things.
 - **Thread:** The "worker" inside the container. Every process has at least one thread (the main thread).
- **Shared Resources:** All threads within a single process share the same memory (heap) and files, but they have their own execution state.
- **The Thread Object:** In the Windows kernel, each thread is represented by an ETHREAD (Executive Thread) block and a KTHREAD (Kernel Thread) block.

Multitasking (Windows)

- **The Thread Object:** In the Windows kernel, each thread is represented by an **ETHREAD** (Executive Thread) block and a **KTHREAD** (Kernel Thread) block.
- **What a Thread Owns:**
 - **Stack:** Two separate stacks—one for User Mode and one for Kernel Mode (to prevent security leaks).
 - **Thread Local Storage (TLS):** Private memory unique to that thread.
 - **Context:** The current state of CPU registers (EIP/RIP, Stack Pointer, etc.).
 - **TEB (Thread Environment Block):** A data structure in user mode used by the thread to manage its own state.

Multitasking (Windows)

- **Multitasking:** The Windows scheduler doesn't wait for a thread to stop. It forcibly interrupts threads when their time is up.
- **Priority Levels (0–31):**
 - **0:** Idle (System only).
 - **1–15:** Dynamic Priorities (Standard apps).
 - **16–31:** Real-Time Priorities (Critical OS tasks, audio processing).
- **Priority Boosting:** Windows automatically "boosts" the priority of a thread if it's in the foreground or if it has just finished waiting for I/O, making the system feel more responsive.

Multitasking (Windows)

- **The Race Condition:** Since threads share the same memory, two threads trying to update the same variable at the same time can cause data corruption.
- **Windows Synchronization Primitives:**
 - **Mutex (Mutual Exclusion):** Only one thread can "own" a resource at a time.
 - **Semaphores:** Allows a specific number of threads (e.g., 5) to access a resource.
 - **Critical Sections:** A high-speed, user-mode locking mechanism for threads within the *same* process.
 - **Events:** A way for one thread to "signal" another that a task is finished.

Shell Programs

- A *shell* is an application program that runs programs on behalf of the user.
 - **sh** Original Unix shell (Stephen Bourne, AT&T Bell Labs, 1977)
 - **csh** BSD Unix C shell (**tcsh**: enhanced **csh** at CMU and elsewhere)
 - **bash** “Bourne-Again” Shell

```
int main() {
    char cmdline[MAXLINE];

    while (1) {
        /* read */
        printf("> ");
        Fgets(cmdline, MAXLINE, stdin);
        if (feof(stdin))
            exit(0);

        /* evaluate */
        eval(cmdline);
    }
}
```

*Execution is a sequence of
read/evaluate steps*

Windows
cmd.exe

What Is a “Background Job”?

- **Users generally run one command at a time**
 - Type command, read output, type another command
- **Some programs run “for a long time”**
 - Example: “delete this file in two hours”

```
unix> sleep 7200; rm /tmp/junk # shell stuck for 2 hours
```

- **A “background” job is a process we don't want to wait for**

```
unix> (sleep 7200 ; rm /tmp/junk) &  
[1] 907  
unix> # ready for next command
```

- **A background process is an executable that runs without direct user interaction or a graphical user interface (GUI).** They handle system-level tasks (disk indexing, network monitoring, security). Unlike foreground apps, they do not have a window and often start automatically during the boot sequence.

Background Processes

- **Service Control Manager (SCM):** A special system process (**services.exe**) that acts as the "boss" for background tasks.
 - Can run even when no user is logged in.
 - Often run under specialized accounts like LocalSystem, NetworkService, or LocalService to minimize security risks.
- In older OS Windows designs, every service was its own .exe. This wasted memory because of redundant overhead.
- **Solution:** svchost.exe is a generic host process name for services that run from dynamic-link libraries (DLLs). Windows groups related services into a single "host" process to save RAM.
- **Time Slicing:** Background processes share the same CPU as foreground apps. The Windows Kernel uses Preemptive Multitasking to give each process a "quantum" (time slice).

User Mode vs. Kernel Mode

- **User Mode (Ring 3):** Most background processes run here. If they crash, they don't take down the whole system. They have restricted access to hardware.
- **Kernel Mode (Ring 0):** Reserved for critical drivers and the core OS. If a background process (like a low-level driver) running here fails, it results in a system crash.
- **System Calls:** Background processes must use *System Calls* to request resources (like writing to a log file on the SSD) from the Kernel.